



Networking with Kubernetes!

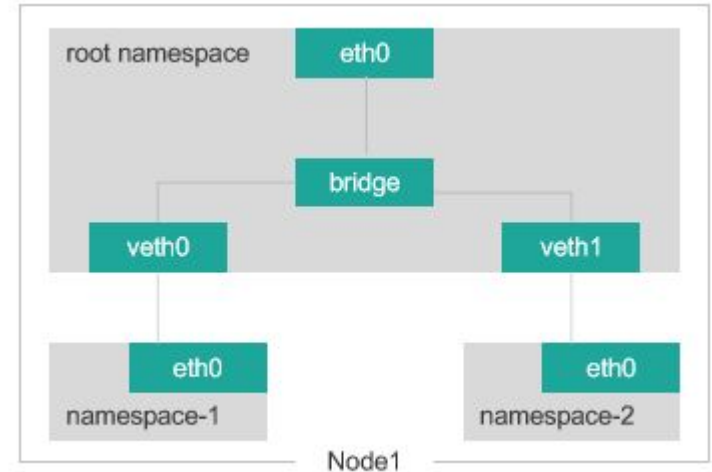


Part 1

Intro: K8S Networking

Namespace

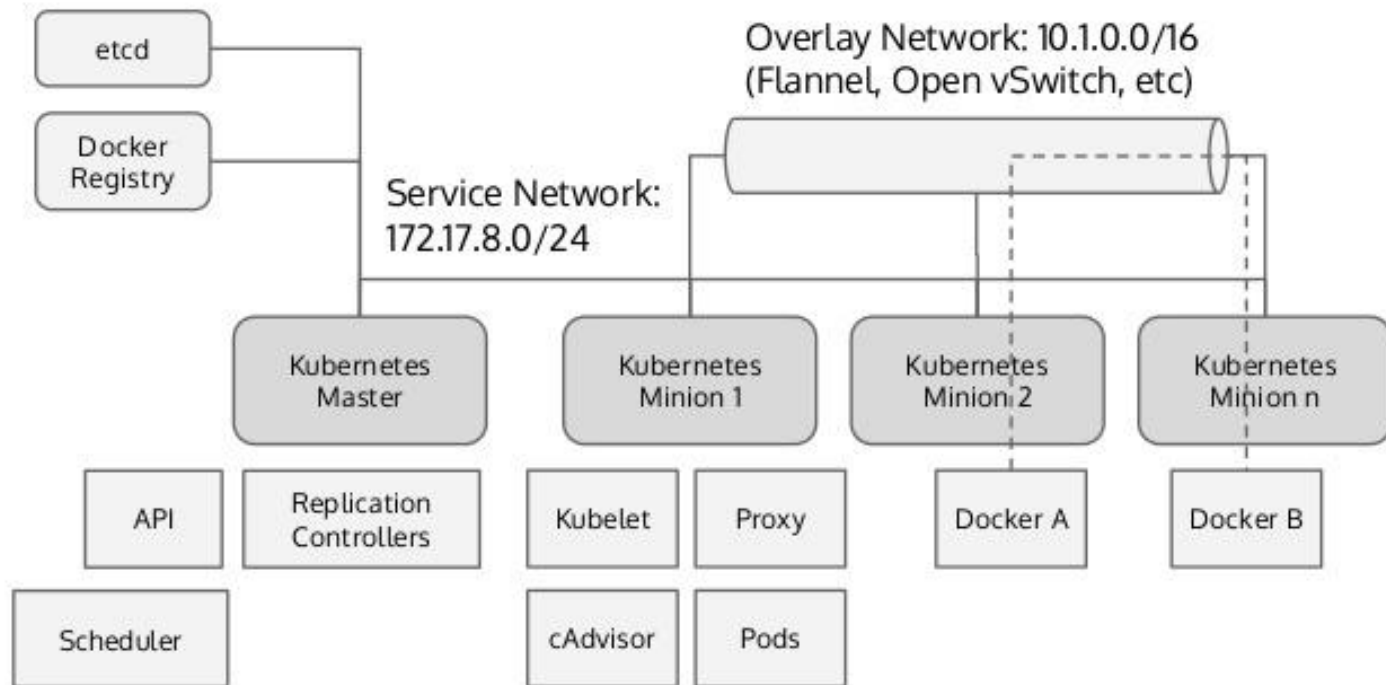
- Linux kernel has 6 types of namespaces: *pid,net,mnt,uts,ipc,user*
- Network namespaces provide a brand-new network stack for all the processes within the namespace.
- That includes network interfaces, routing tables and iptables rules



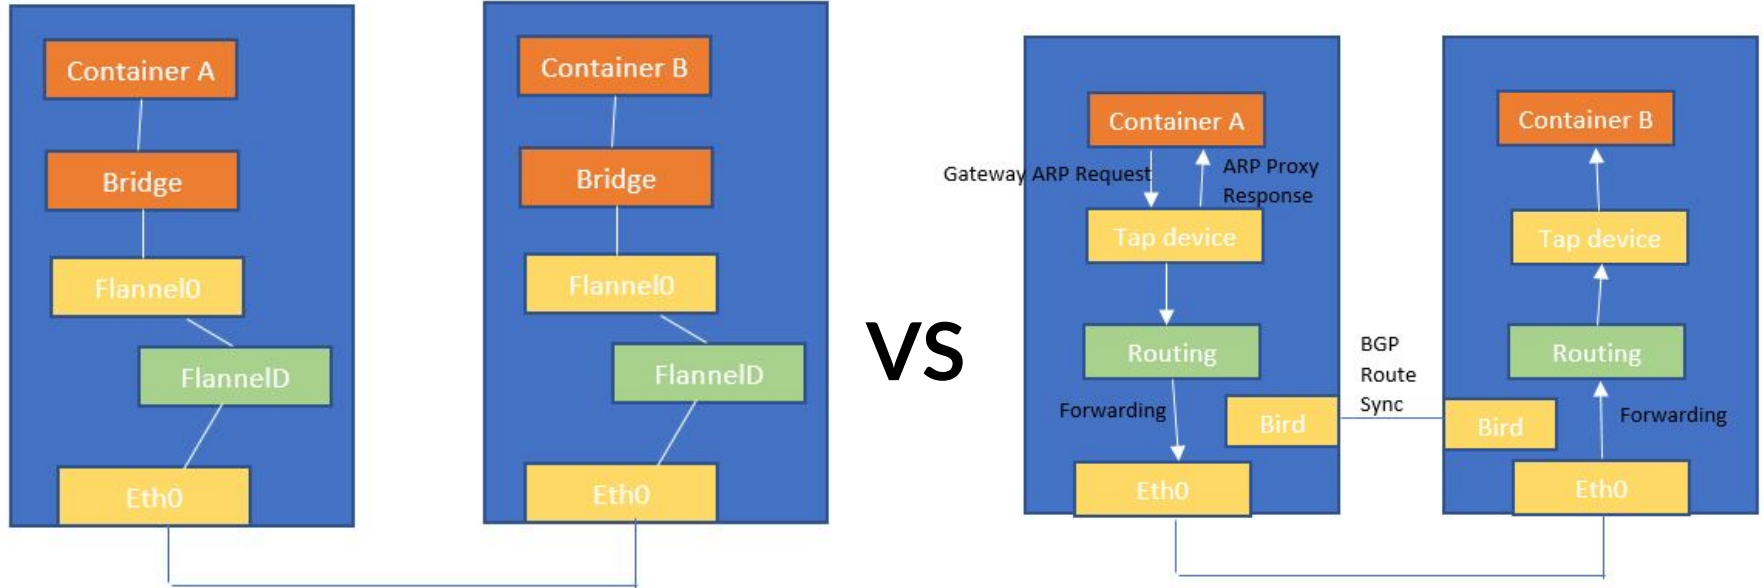
Live demo: Play with namespaces and cgroups.



Kubernetes Architecture



Flannel VS Calico



Pods

Pods

- Lowest common denominator in K8S. Pod is comprised of one or more containers along with a "pause" container
- Pause container act as the "parent" container for other containers inside the pod. One of it's primary responsibilities is to bring up the network namespace
- Great for the redundancy: Termination of other containers do not result in termination of the network namespace

```
[root@Master1 ~]# kubectl get pods --namespace web -o wide
```

NAME	READY	STATUS	AGE	IP	NODE
nginx-deployment-76bf4969df-5xzv7	1/1	Running	7m53s	172.31.155.17	Worker-1

↓

```
[root@Worker-1 ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	NAMES
93490bfce728	docker.io/nginx@sha256	"nginx -g 'daemon off'"	47 seconds ago	k8s_nginx_nginx-deployment
2ef012ea5db0	k8s.gcr.io/pause:3.1	"/pause"	57 seconds ago	k8s_POD_nginx-deployment

Accessing Pod Namespaces

Accessing Pod Namespaces

- Multiple ways to access pod namespaces
- 'kubectl exec -it'
- 'docker exec -it'
- nsenter ("namespace enter", let you run commands that are installed on the host but not on the container)

```
[root@worker-1 ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	NAMES
5b54f2a44c3b	d8233ab899d4	"sleep 3600"	35 minutes ago	k8s_busybox_busybox0-6hc7c
43e42c45522b	k8s.gcr.io/pause:3.1	"/pause"	10 hours ago	k8s_POD_busybox0-6hc7c

```
[root@worker-1 ~]# docker inspect -f '{{.State.Pid}}' 5b54f2a44c3b
21388

[root@worker-1 ~]# nsenter -t 21388 -n ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
17: eth0@if18: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP
    link/ether fa:4d:26:0b:4a:c7 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 192.168.1.196/32 brd 192.168.1.196 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::f84d:26ff:fe0b:4ac7/64 scope link
        valid_lft forever preferred_lft forever
```

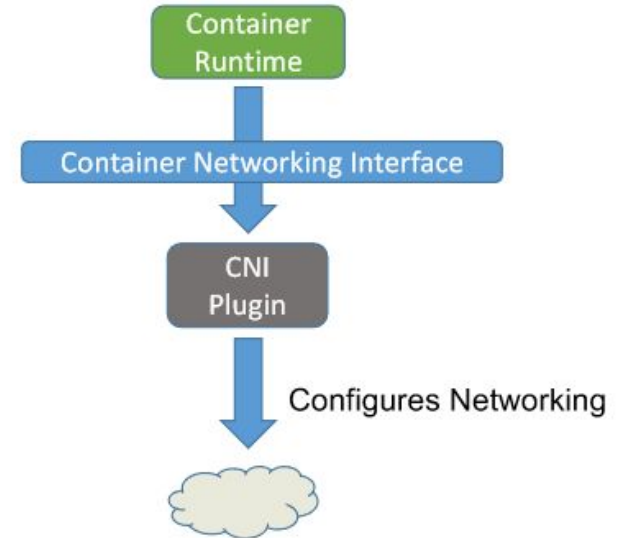
```
[root@worker-1 ~]# docker inspect -f '{{.State.Pid}}' 43e42c45522b
8112

[root@worker-1 ~]# nsenter -t 8112 -n ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
17: eth0@if18: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP
    link/ether fa:4d:26:0b:4a:c7 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 192.168.1.196/32 brd 192.168.1.196 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::f84d:26ff:fe0b:4ac7/64 scope link
        valid_lft forever preferred_lft forever
```

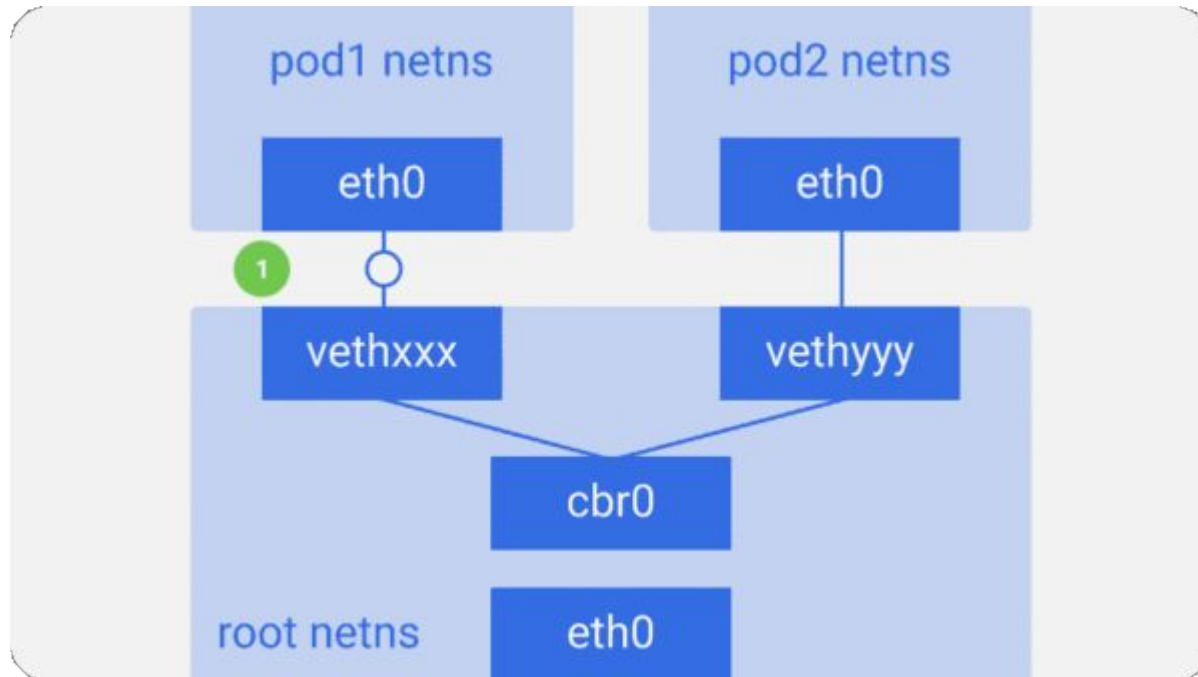
Both containers belong to the same pod => Same Network Namespace => same 'ip a' output

Container Networking Interface : CNI

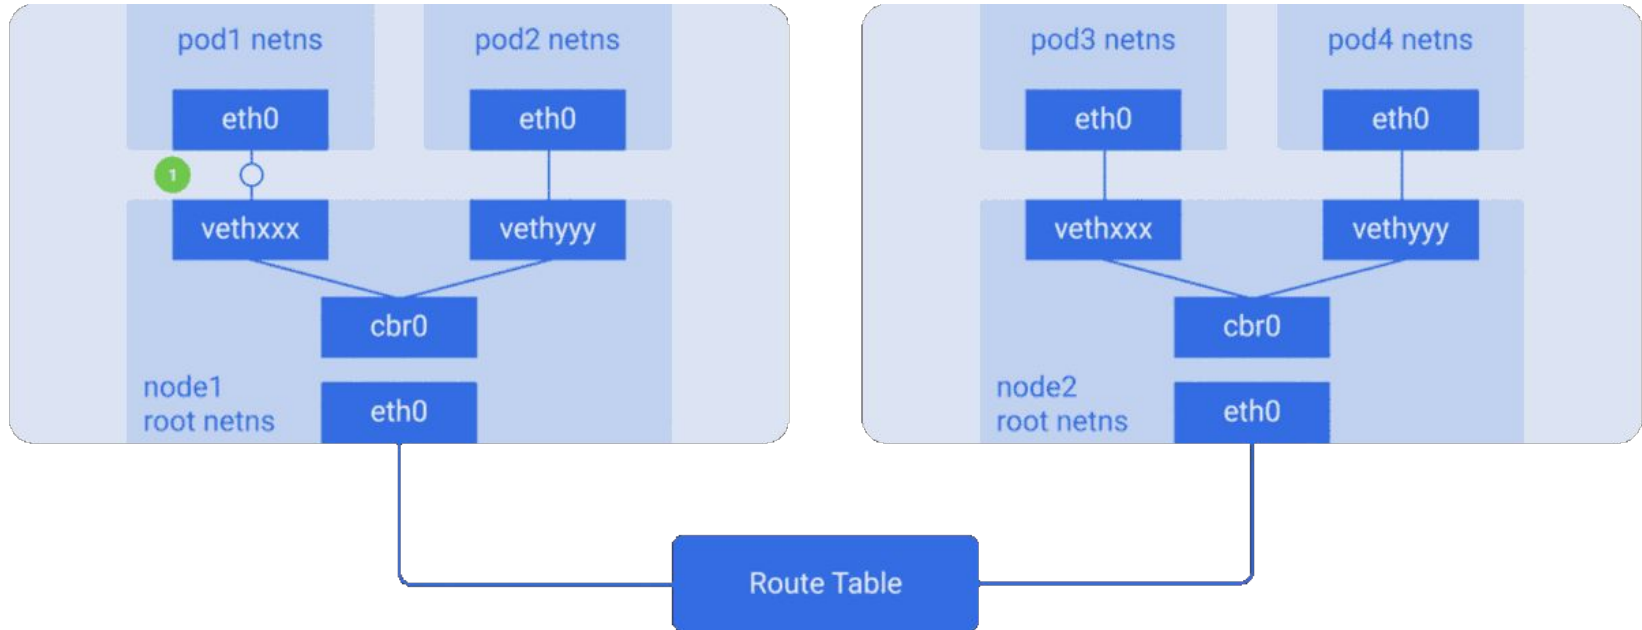
- Interface between container runtime and network implementation
- Network plugin implements the CNI spec. It takes a container runtime and configure (attach/detach) it to the network
- CNI plugin is an executable (in: /opt/cni/bin)
- When invoked it reads in a JSON config & Environment Variables to get all the required parameters to configure the container with the network



Intra-node communication



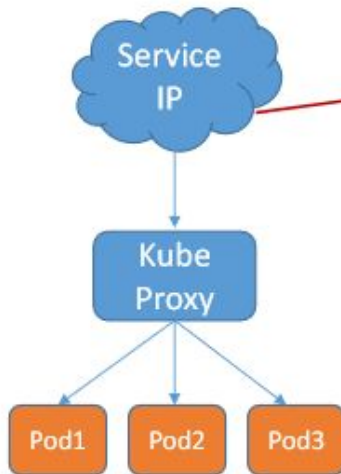
Inter-node communication



Back to the Basics

Services

- Pods are mortal
- Need a higher level abstractions: Services
- “Service” in Kubernetes is a conceptual concept. Service is not a process/daemon. Outside networks doesn't learn Service IP addresses
- Implemented through Kube Proxy with IPTables rules



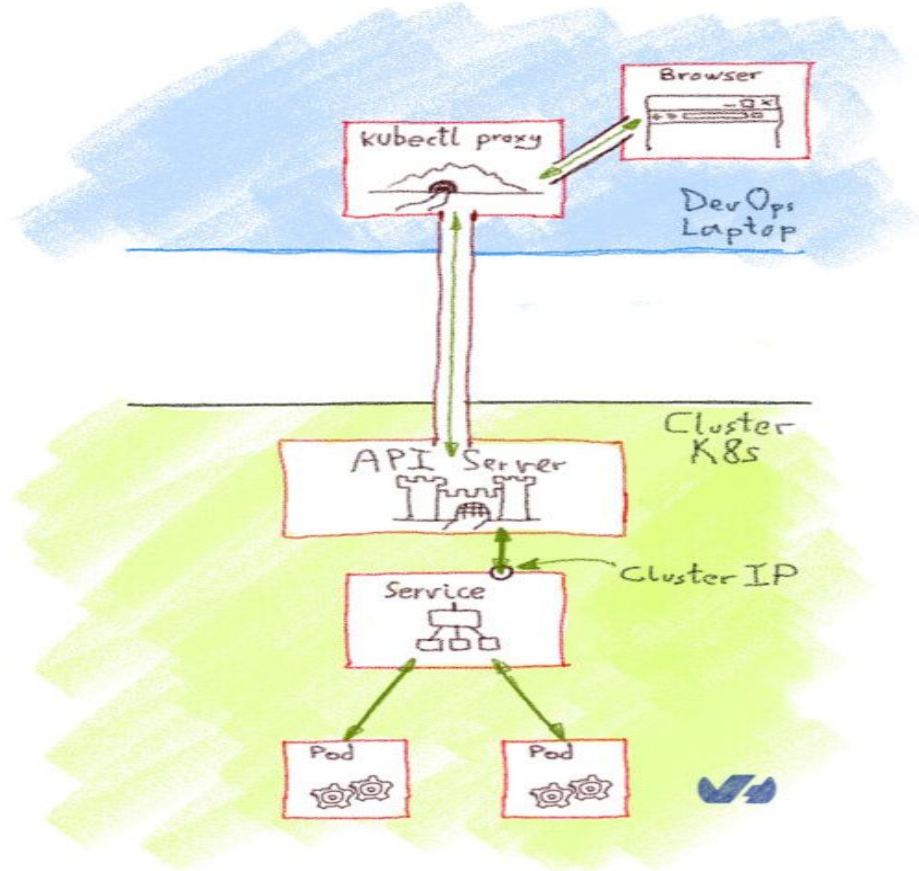
```
$ kubectl get services -n demo -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
hostnames	ClusterIP	10.96.13.117	<none>	80/TCP	23h	app=hostnames

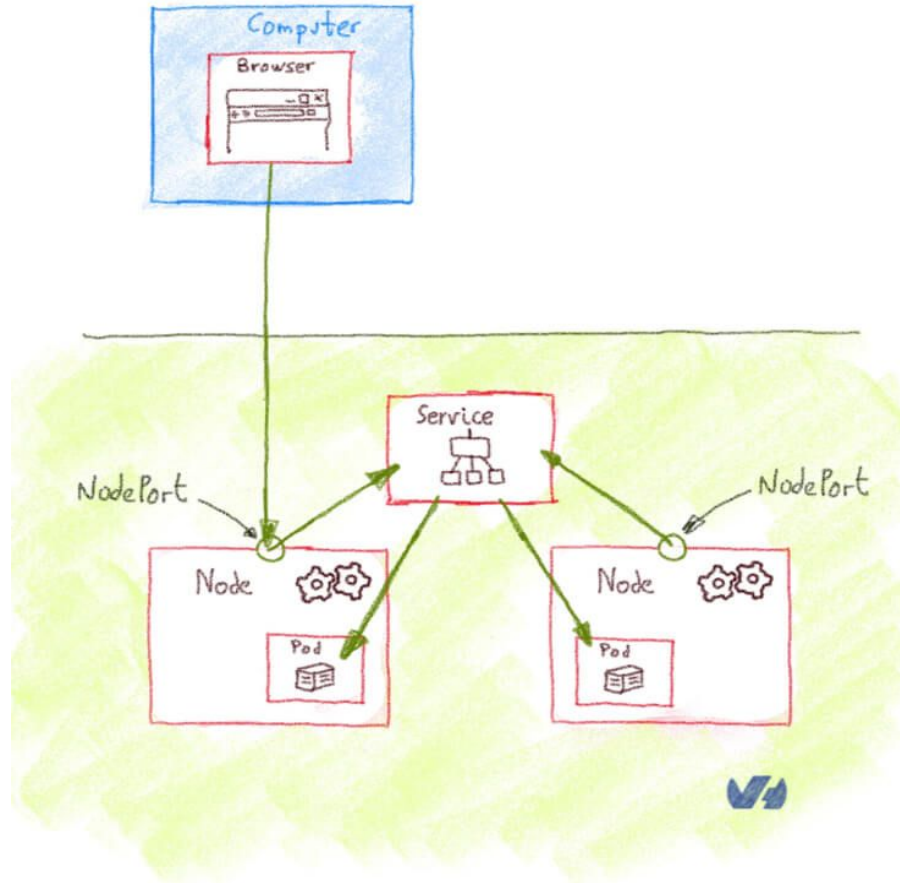
```
$ kubectl describe service -n demo
```

Name: hostnames
Namespace: demo
Labels: <none>
Annotations: <none>
Selector: app=hostnames
Type: ClusterIP
IP: 10.96.13.117
Port: default 80/TCP
TargetPort: 9376/TCP
Endpoints: 192.168.1.63:9376,192.168.2.171:9376,192.168.3.155:9376
Session Affinity: None
Events: <none>

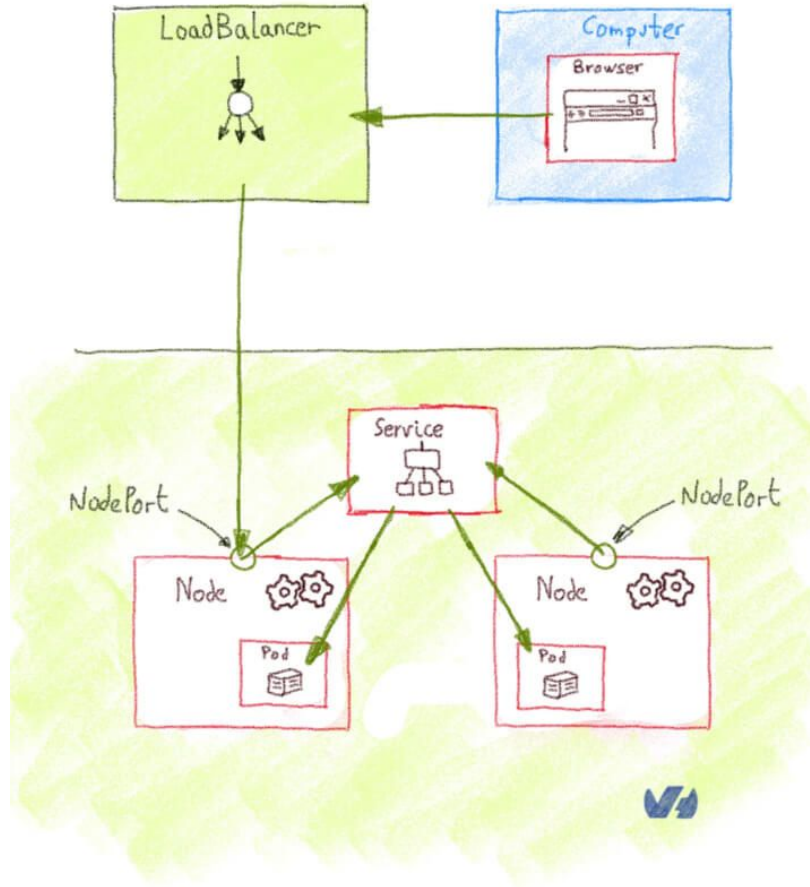
ClusterIP: Service is accessed via 'ClusterIP'



NodePort: Service is accessed via 'NodeIP:port'

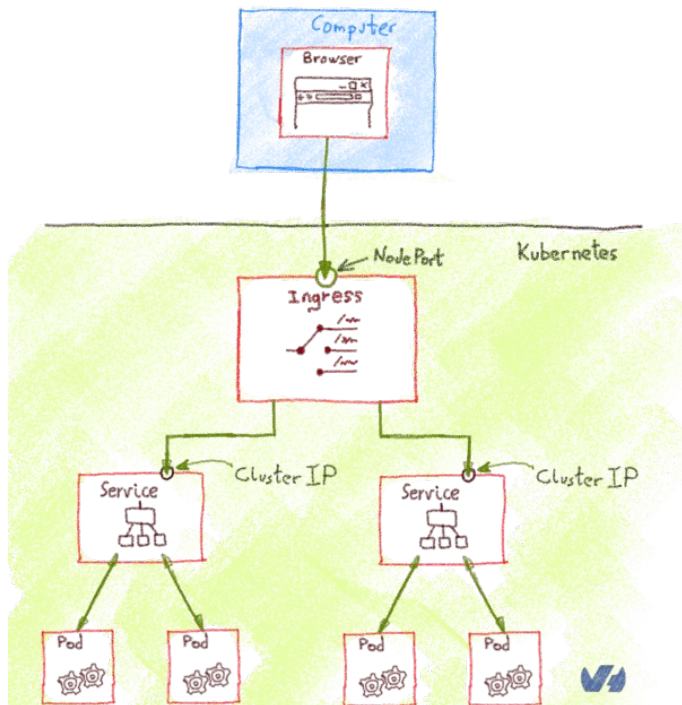


LoadBalancer: Service is accessed via Loadbalancer

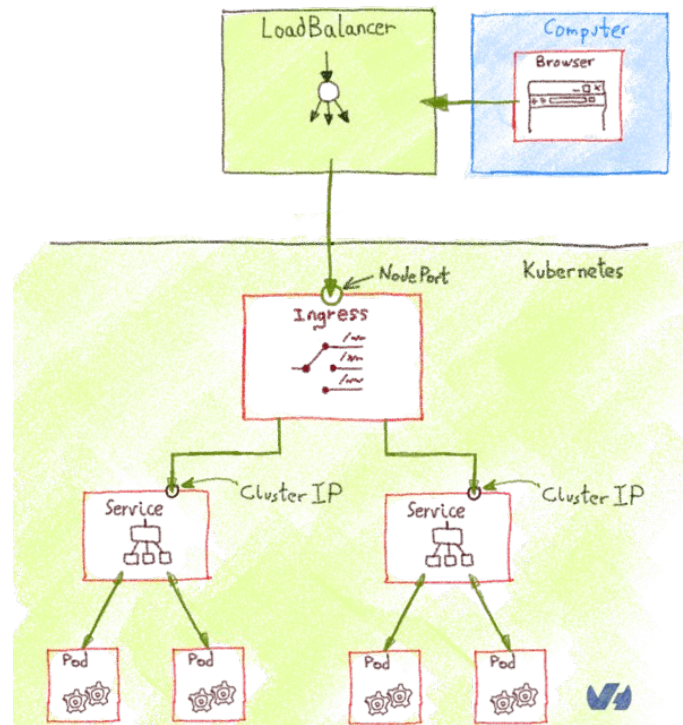


Ingress

Ingress is an API object that manages external access to the services in a cluster.



Ingress



Ingress behind Load Balancer



Part 2:

Setting up K8S The Hard Way

On top of Amazon Web Services (AWS)



1- Provisioning Compute Resources

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/03-compute-resources.md>

→ **Networking:**

- VPC
- Subnet
- Internet Gateway
- Route Tables
- Security Groups (aka Firewall Rules)

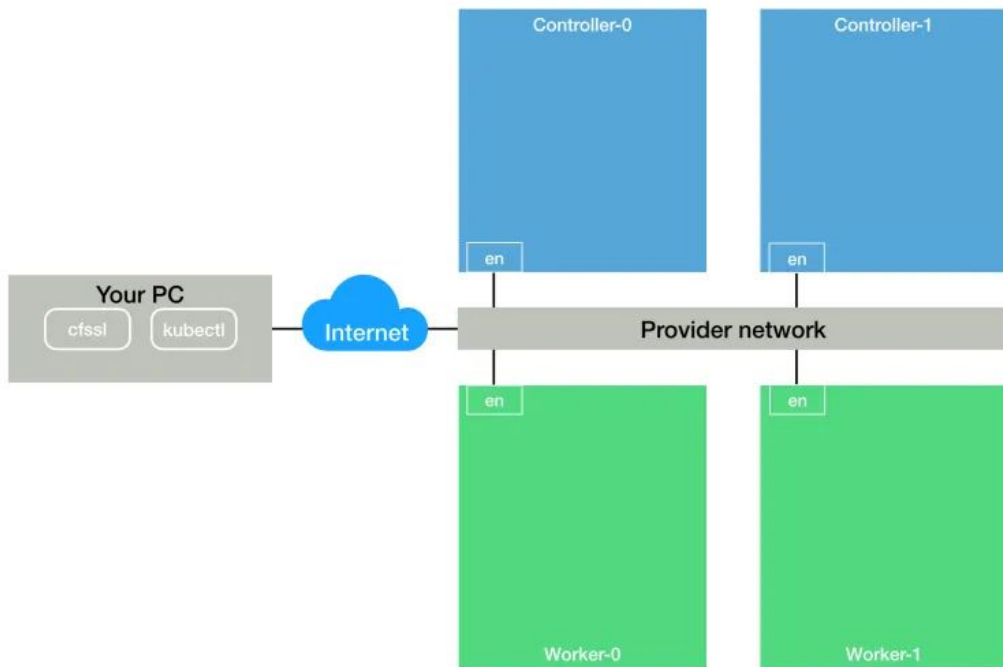
→ Create a Network Load Balancer

→ **Compute Instances:**

- Instance Image + SSH Key Pair
- Kubernetes Controllers
- Kubernetes Workers

1- Provisioning Compute Resources

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/03-compute-resources.md>





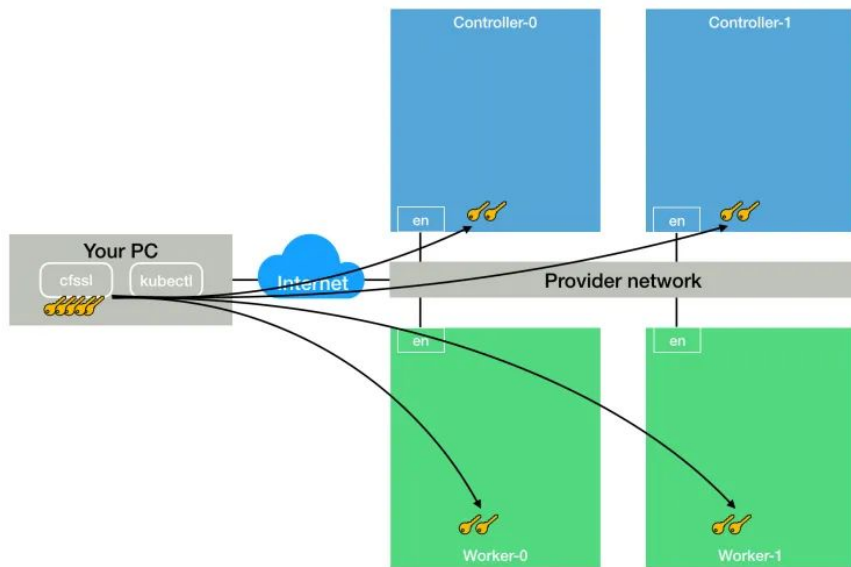
2- Provisioning a CA and Generating TLS Certificates

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/04-certificate-authority.md>

- **Certificate Authority**
- **Client and Server Certificates:**
 - The Admin Client Certificate
 - The Kubelet Client Certificates
 - The Controller Manager Client Certificate
 - The Kube Proxy Client Certificate
 - The Scheduler Client Certificate
 - The Kubernetes API Server Certificate
- **The Service Account Key Pair**
- **Distribute the Client and Server Certificates**

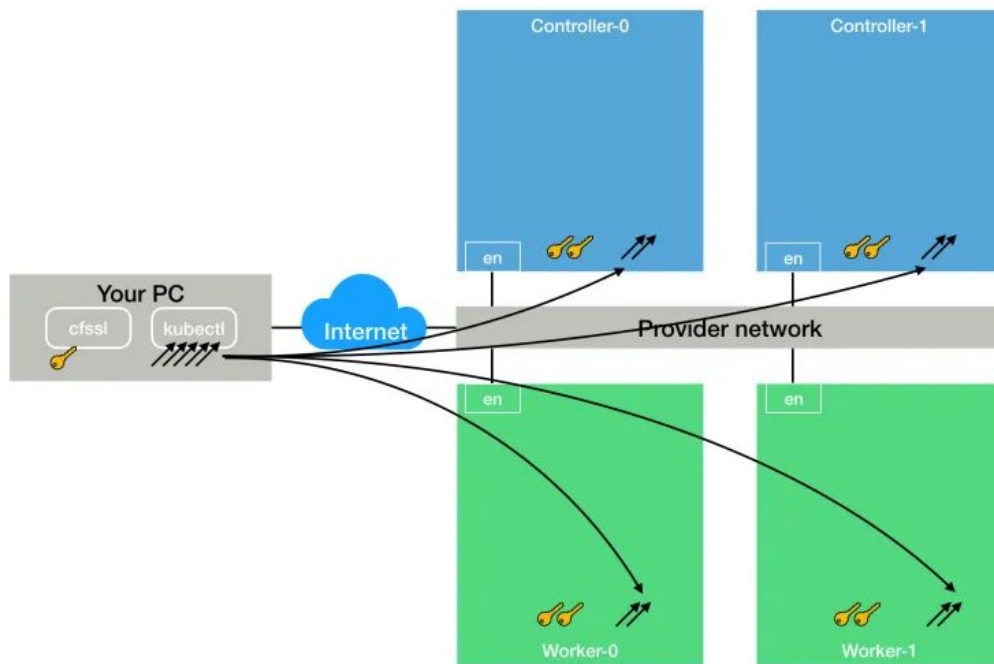
2- Provisioning a CA and Generating TLS Certificates

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/04-certificate-authority.md>



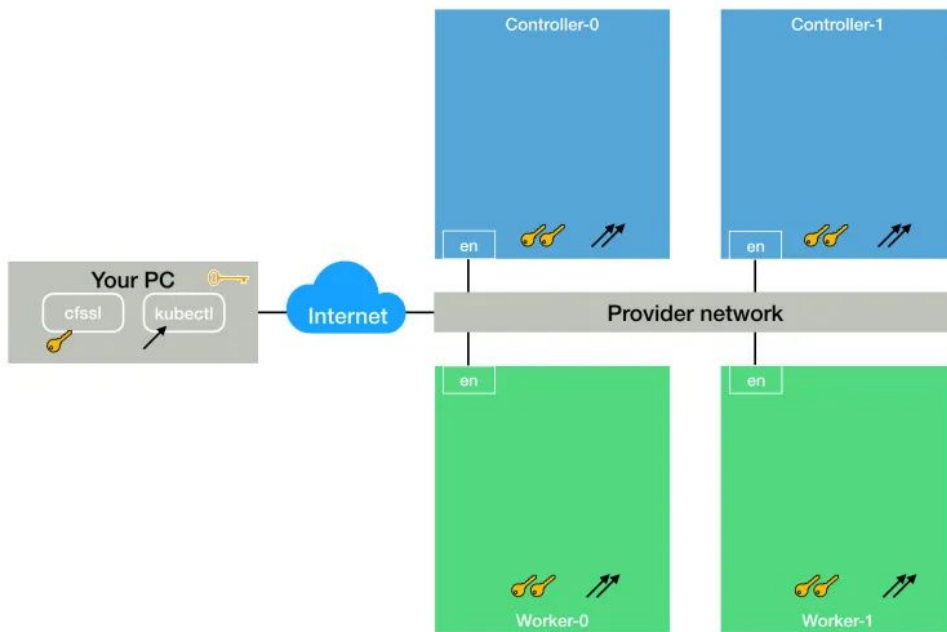
3- Generating Kubernetes Configuration Files for Authentication

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/05-kubernetes-configuration-files.md>



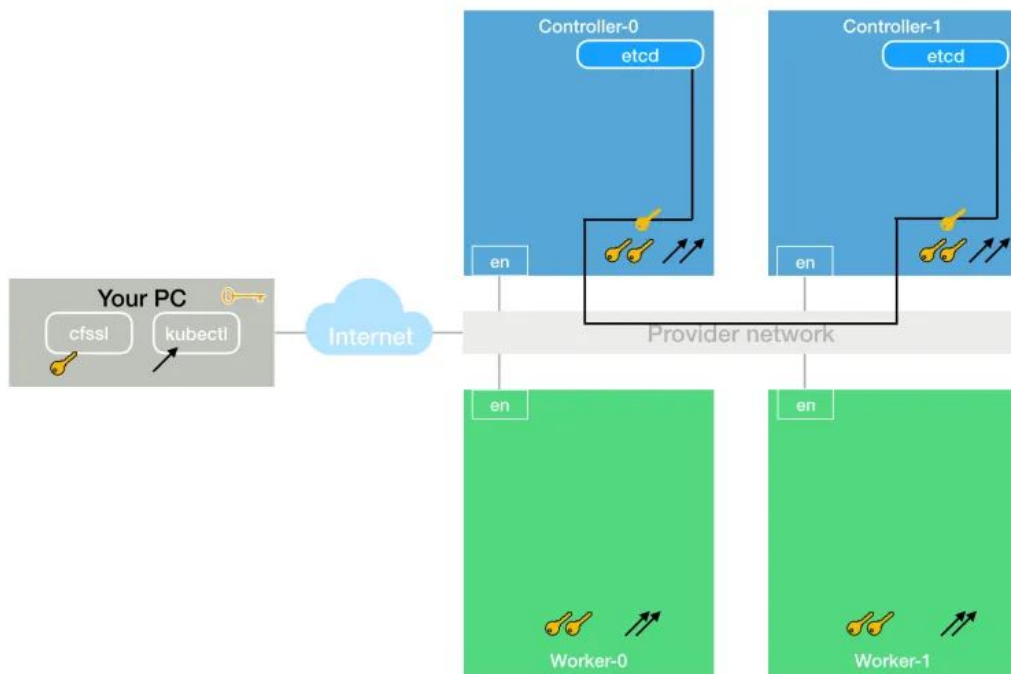
4- Generating the Data Encryption Config and Key

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/06-data-encryption-keys.md>



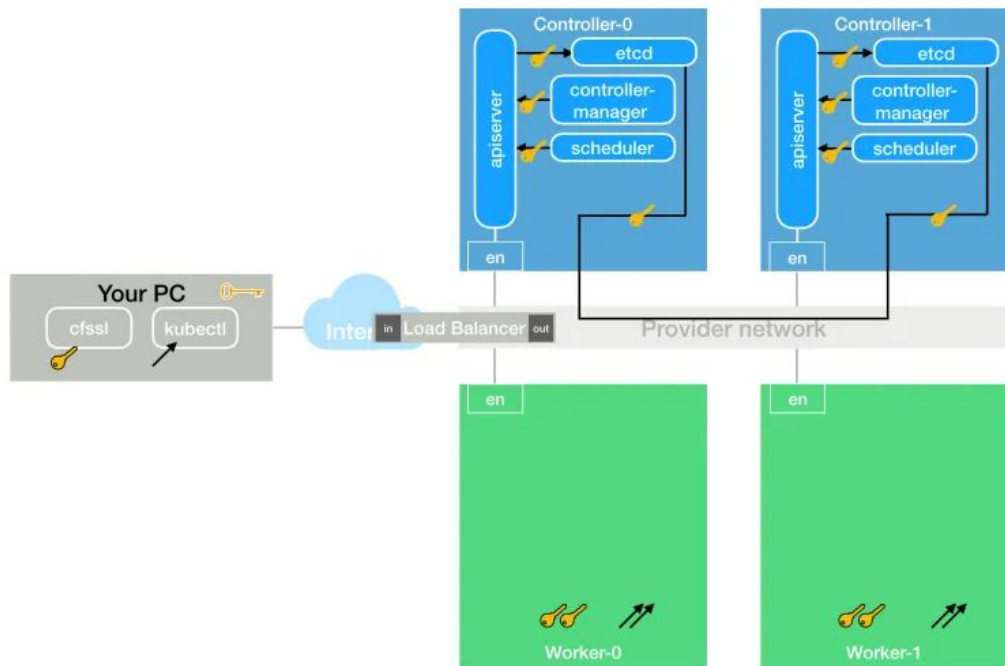
5- Bootstrapping the etcd Cluster

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/07-bootstrapping-etcd.md>



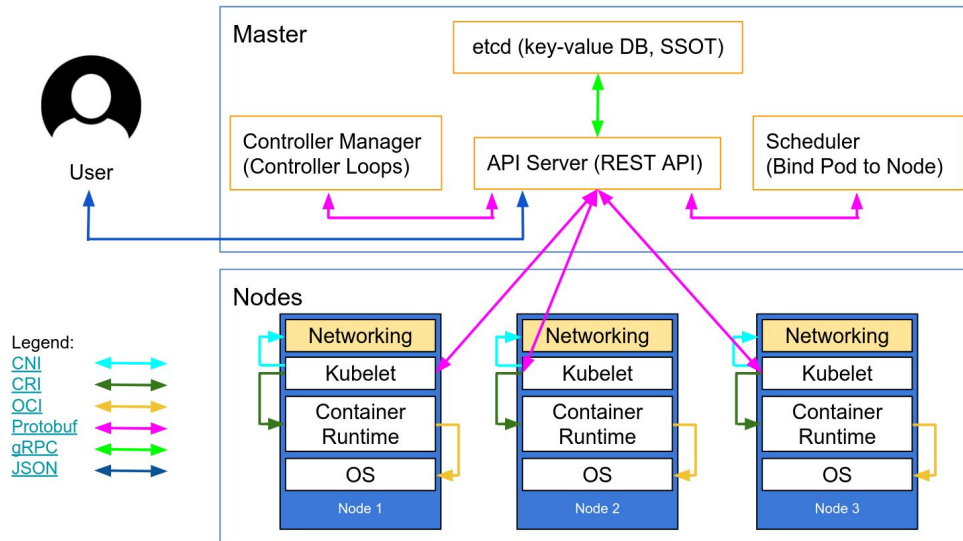
6- Bootstrapping the Kubernetes Control Plane

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/08-bootstrapping-kubernetes-controllers.md>



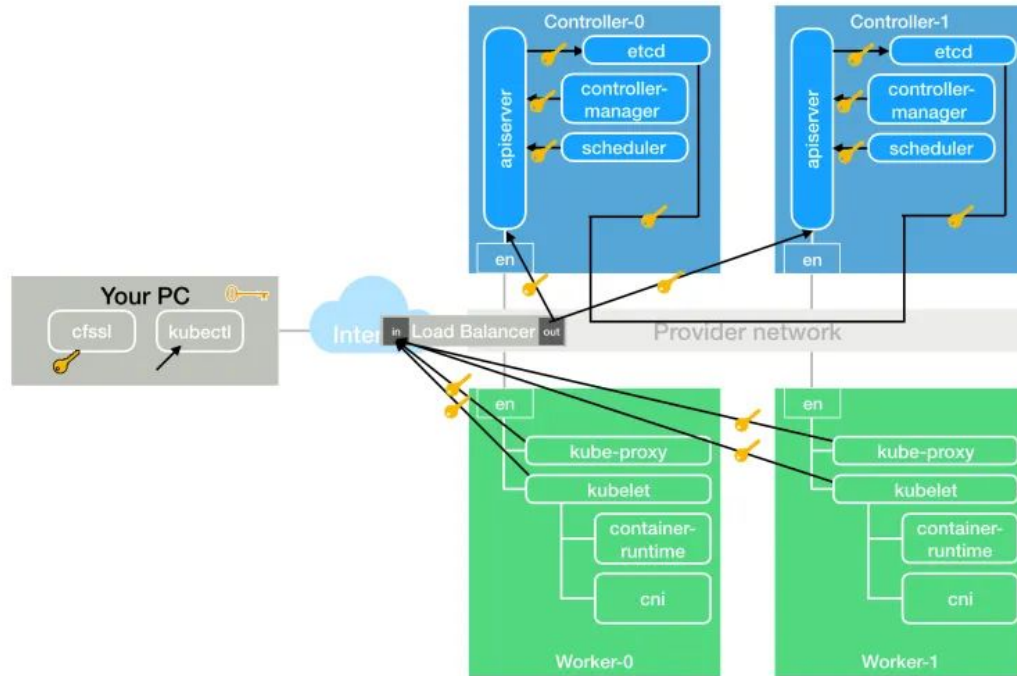
6- Bootstrapping the Kubernetes Control Plane

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/08-bootstrapping-kubernetes-controllers.md>



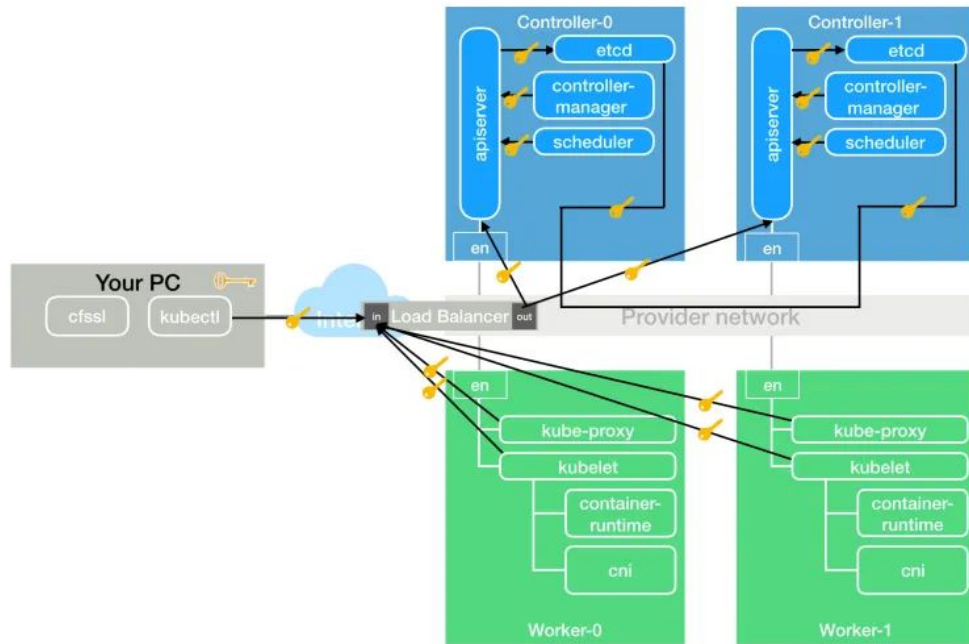
7- Bootstrapping the Kubernetes Worker Nodes

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/09-bootstrapping-kubernetes-workers.md>




8- Configuring kubectl for Remote Access

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/10-configuring-kubectl.md>



9- Provisioning Pod Network Routes

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/11-pod-network-routes.md>

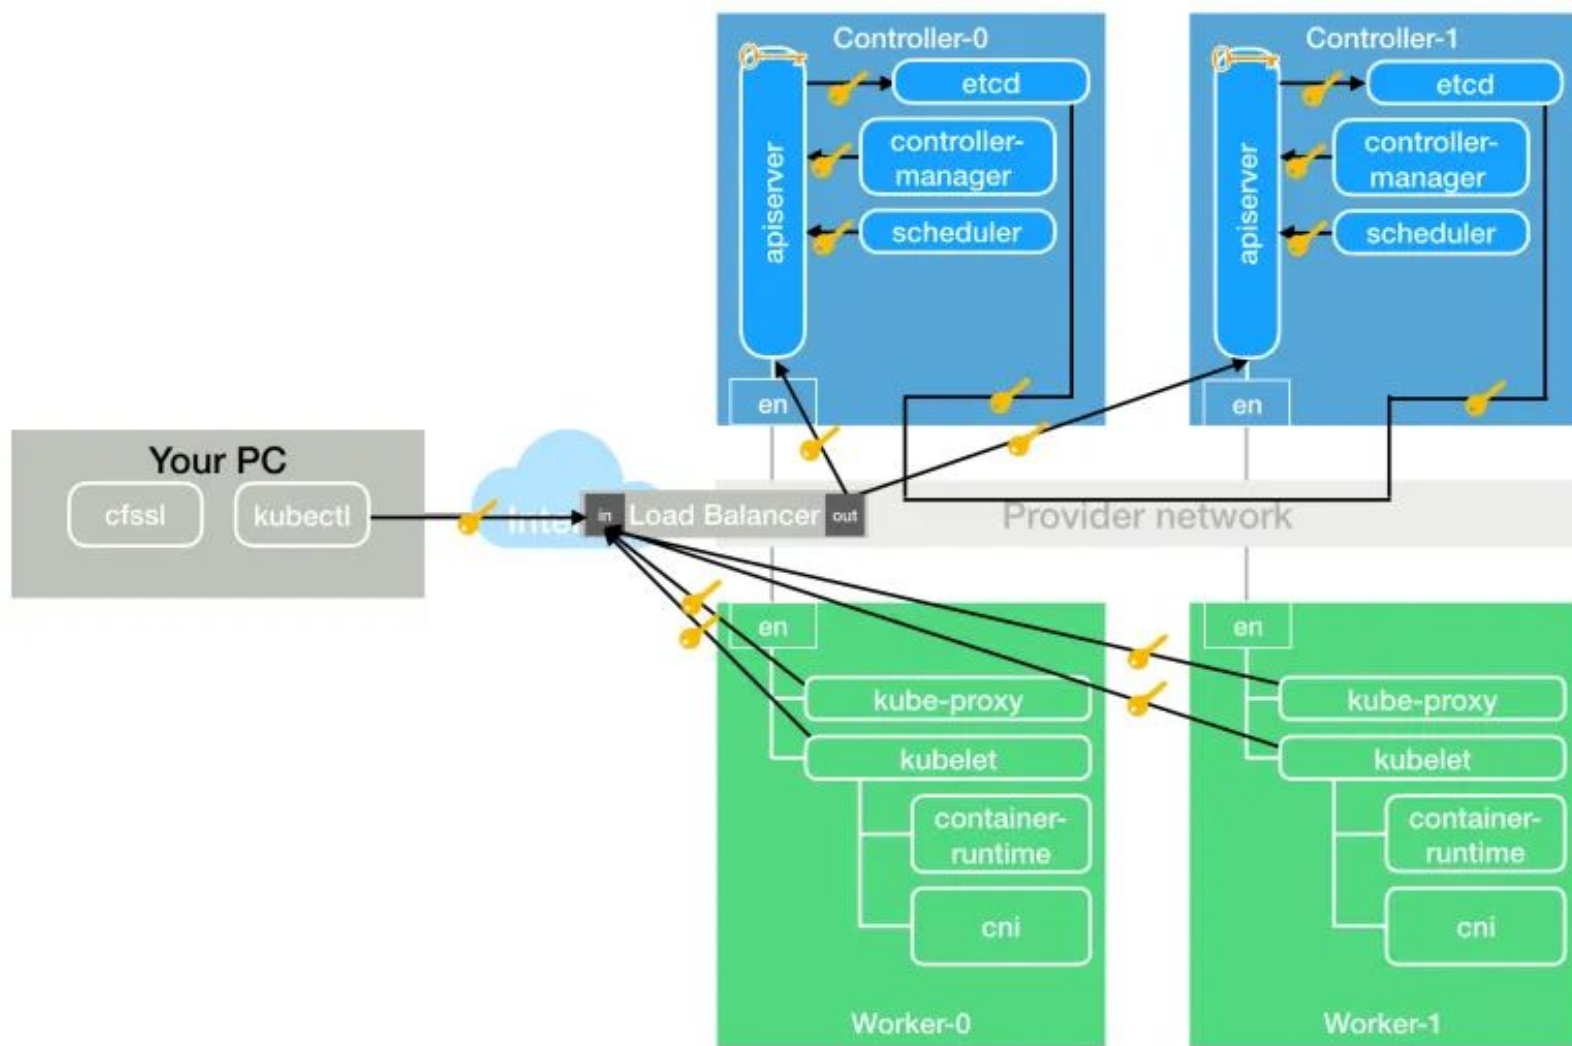




10- Deploying the DNS Cluster Add-on

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/12-dns-addon.md>

1





12- Cleaning Up

<https://github.com/TunisJAM/kubernetes-the-hard-way-aws/blob/master/docs/14-cleanup.md>



Part 4:

Istio, a modern service mesh

ISTIO architecture

